

4 sawal

1. Suggest some AI innovations ideas for this which u said are missing also explain them thoroughly along with this compare the existing approaches {both} in detail.
2. Compare current implementaion decisions wrt pipeline recommendations and suggest the better of all, depending upon usecases.
3. As per the face verification part, suggest the methodologies to smoothly implement it as per the real time verification modules just like mediapipe or YOLO on camerafeed itself.
4. Validate whether deepface could be appropriate for facial recognition in real time on camera feed/pictures

+ Thought: 5.2s

Technical Analysis & Innovation Roadmap

1. AI INNOVATIONS: Missing Components — Deep Dive & Comparison

1.1 PRNU (Photo-Response Non-Uniformity) Noise Residual Analysis

What it is: Every CMOS/CCD sensor has a unique, fixed-pattern noise "fingerprint" (PRNU) caused by manufacturing variations in pixel sensitivity. When a region is spliced from another image (different camera), its PRNU fingerprint differs from the host image.

How it works:

1. Extract noise residual: $W = I - \text{denoise}(I)$ (wavelet/BM3D denoising)
2. Estimate camera fingerprint: $K = \text{average}(W_i)$ over many flat-field images from same device
3. Cross-correlate: $p = \text{corr}(W_{\text{test}}, K_{\text{reference}})$ per block (e.g., 256×256)
4. Regions with $p \ll \text{global average}$ = foreign sensor = splice

Comparison vs Current ELA:

Dimension	ELA (Current)	PRNU (Proposed)
-----------	---------------	-----------------

Detects	Re-compression artifacts (JPEG quantization differences)	Sensor-level splicing (different camera/phone)
---------	--	--

False Positives	High on re-scans, WhatsApp forwards, screenshots	Low — PRNU survives recompression, scaling, mild filtering
-----------------	--	--

False Negatives	Low for crude copy-paste at same quality	High if forged region from same device, or heavy post-processing
-----------------	--	--

Data Needs None (single image) Reference images from same device (or blind estimation via Expectation-Maximization)

Compute ~50ms (1 recompress + diff) ~500ms–2s (denoising + correlation per block)

SIH Relevance Catches "photo swapped in Photoshop" Catches "photo printed from another phone + re-photographed"

Implementation Path:

Minimal viable PRNU (blind, no reference)

from scipy.ndimage import gaussian_filter

import pywt

```
def extract_prnu_residual(image, sigma=2.0):
```

```
    """Wavelet-based noise residual (simplified)."""
```

```
    coeffs = pywt.wavedec2(image, 'db1', level=3)
```

```
    # Zero approximation coefficients, reconstruct detail-only
```

```
    coeffs[0] *= 0
```

```
    denoised = pywt.waverec2(coeffs, 'db1')
```

```
    return image - denoised
```

```
def prnu_correlation_map(residual, block_size=256):
```

```
    """Sliding window normalized cross-correlation."""
```

```
    h, w = residual.shape
```

```
    corr_map = np.zeros((h//block_size, w//block_size))
```

```
    global_fp = residual.mean(axis=(0,1)) # crude global fingerprint
```

```
    for i in range(0, h, block_size):
```

```
        for j in range(0, w, block_size):
```

```
            block = residual[i:i+block_size, j:j+block_size]
```

```
            if block.size == 0: continue
```

```
            corr_map[i//block_size, j//block_size] = np.corrcoef(block.flatten(),  
global_fp.flatten())[0,1]
```

```
return corr_map
```

Verdict: High impact for SIH. PRNU is the only technique that detects cross-device photo substitution (printed photo, screen replay) — the exact "altered photographs" threat in the problem statement. ELA alone cannot.

1.2 Stamp/Seal Template Matching (Targeted Forgery Detection)

What it is: Visa stamps, holograms, embassy seals have known templates. Forge = stamp cloned from another visa, or stamp pixels manipulated.

Approaches:

Approach	Description	Pros	Cons
SIFT/ORB + Homography	Rotation/scale invariant	Match stamp template → detect if transformed copy	Fails if stamp heavily degraded
Siamese CNN (Patch Verification)	Learns subtle variations (ink bleed, pressure)	Train $f(\text{patch}) \rightarrow \text{embedding}$; genuine stamps cluster together	Needs labeled genuine stamps per visa type
Frequency-Domain Watermark Check	Invisible, robust to printing	Many official stamps embed DCT-domain watermarks	Requires issuer cooperation

Comparison vs Current Copy-Move:

Dimension	Copy-Move (Current)	Stamp Template Matching
Detects	Any cloned region (generic)	Specifically: forged visa stamps, seals
Threat Model	"Clone tool used anywhere"	"Stamp copied from another visa" (explicit in problem stmt)
False Positives	High on repetitive textures (holograms, patterns)	Low — only triggers on stamp zones

Implementation 50 lines ORB 200+ lines (template registry + alignment + scoring)

Implementation Path:

Stamp zone localization (rule-based from MRZ/OCR)

```
def locate_stamp_zone(doc_type, image, mrz_bbox):
```

```
    if doc_type == "visa":
```

```
        # Visa stamps typically bottom-right, below MRZ
```

```
        h, w = image.shape[:2]
```

```
        return image[mrz_bbox[1]+mrz_bbox[3]:, w//2:] # heuristic
```

```
return None
```

```
# Template matching with ORB + RANSAC
```

```
def verify_stamp_authenticity(stamp_region, template_db):
```

```
    best_score = 0
```

```
    for visa_type, templates in template_db.items():
```

```
        for tmpl in templates:
```

```
            matches = orb_match(stamp_region, tmpl)
```

```
            if len(matches) > MIN_MATCHES:
```

```
                H, mask = cv2.findHomography(...)
```

```
                inlier_ratio = mask.sum() / len(matches)
```

```
                best_score = max(best_score, inlier_ratio)
```

```
    return best_score > THRESHOLD
```

Verdict: Directly addresses "tampered visa stamps" in problem statement. Current copy-move is too generic; this is targeted, explainable, and officer-verifiable.

1.3 LayoutLM / Document Understanding Transformer (Donut) for Field Extraction

What it is: Multimodal transformers that jointly encode text + layout (bbox) + image for semantic field extraction.

Comparison vs Current Regex/Heuristic Extractors:

Dimension	Regex/Heuristic (Current)	LayoutLMv3 / Donut
-----------	---------------------------	--------------------

Accuracy on Clean Docs	~85%	~95%+
------------------------	------	-------

Accuracy on Noisy/Rotated	~40%	~80%+
---------------------------	------	-------

New Document Types	Write new regex per type	Fine-tune on 50–100 samples
--------------------	--------------------------	-----------------------------

Explainability	High (readable rules)	Medium (attention maps)
----------------	-----------------------	-------------------------

Compute	~5ms	~200ms (CPU) / ~30ms (GPU)
---------	------	----------------------------

Dependencies	None	transformers, torch, model weights (~400MB)
--------------	------	---

Why It Matters for SIH:

- Problem statement: "passports, visas, national IDs, driving licences, permits" — 5+ document types with different layouts

- Regex approach: 5× maintenance burden, brittle on real checkpoint photos
- LayoutLM: Single model, fine-tunable per issuer (Indian passport vs US visa vs Nepal permit)

Implementation Path:

Donut (OCR-free) is lighter for hackathon

```
from transformers import DonutProcessor, VisionEncoderDecoderModel
```

```
processor = DonutProcessor.from_pretrained("naver-clova-ix/donut-base-finetuned-cord-v2")
```

```
model = VisionEncoderDecoderModel.from_pretrained("naver-clova-ix/donut-base-finetuned-cord-v2")
```

```
def extract_fields_donut(image, document_type):
```

```
    # Donut uses task-specific prompts
```

```
    prompt = f"<s_{document_type}>"
```

```
    pixel_values = processor(image, return_tensors="pt").pixel_values
```

```
    decoder_input_ids = processor.tokenizer(prompt, return_tensors="pt").input_ids
```

```
    outputs = model.generate(pixel_values, decoder_input_ids=decoder_input_ids,
max_length=512)
```

```
    return processor.token2json(processor.batch_decode(outputs)[0])
```

Verdict: Replace field_extractors.py entirely. One model handles all document types. Fine-tune on Indian document samples for production.

1.4 ArcFace / AdaFace + Liveness Detection (Face Verification)

What it is: ArcFace (additive angular margin loss) produces 512-d embeddings where cosine distance = identity similarity. SOTA on LFW, CFP-FP, AgeDB.

Liveness Detection (Critical for Border Security):

Method	Type	Detects	Inference
Silent Face Anti-Spoofing (MiniFASNet)			
~10ms		Passive (texture)	Print, screen replay

Eye Blink / Head Pose Challenge Active All presentation attacks ~2s (user interaction)

rPPG (Remote Photoplethysmography) Passive (blood flow) Print, screen, mask ~1s (needs 5+ sec video)

Depth/IR (Hardware) Active All Requires specialized camera

Comparison vs Current Histogram+ORB:

Dimension Histogram+ORB (Current) ArcFace + Liveness

Same Person (Diff Pose/Age) ~0.4–0.6 similarity ~0.7–0.9 cosine sim

Different Person ~0.3–0.5 (high overlap) ~0.1–0.3 (clear separation)

Printed Photo Attack PASSES (texture matches) BLOCKED (liveness fails)

Screen Replay PASSES BLOCKED

Deepfake Video PASSES BLOCKED (rPPG or challenge)

Explainability Low (black-box-ish) Medium (distance + liveness score)

Implementation Path:

insightface (ArcFace) + MiniFASNet (liveness) — both ONNX, CPU-friendly

import insightface

from insightface.app import FaceAnalysis

app = FaceAnalysis(name='buffalo_l', providers=['CPUExecutionProvider'])

app.prepare(ctx_id=0, det_size=(640, 640))

def verify_with_liveness(doc_face_img, live_video_frames):

 # 1. Liveness on live frames (MiniFASNet)

 liveness_scores = [minifasnet.predict(f) for f in live_video_frames]

 if np.mean(liveness_scores) < 0.5:

 return {"match": False, "reason": "Liveness check failed — possible spoof"}

 # 2. ArcFace embedding

 doc_emb = app.get(doc_face_img)[0].embedding

```
live_emb = app.get(live_video_frames[-1])[0].embedding

sim = np.dot(doc_emb, live_emb) / (np.linalg.norm(doc_emb) * np.linalg.norm(live_emb))
```

```
return {"match": sim > 0.45, "similarity": float(sim), "liveness":
float(np.mean(liveness_scores))}
```

Verdict: Non-negotiable for production. Current face module is a security hole. ArcFace + passive liveness is the minimum viable.

1.5 Isolation Forest / Autoencoder for Anomaly Detection (Validation Layer)

What it is: Unsupervised anomaly detection on document field statistics (font size, position, spacing, color histogram) vs known-good templates.

Comparison vs Current Rules Engine:

Dimension	Rules Engine (Current)	Isolation Forest / Autoencoder
Catches Known Violations	✅ Perfect	❌ Not designed for this
Catches Novel Forgeries	❌ Only what rules cover	✅ Flags "looks weird vs template"
Explainability	✅ "Field X failed rule Y" font_size, y_pos"	⚠️ "Anomaly score 0.92 — feature importance: font_size, y_pos"

Training Data None needed 100+ genuine samples per doc type

False Positive Rate 0% (deterministic) 2–5% (tunable)

Hybrid Approach (Recommended):

Rules Engine (hard gates: checksum, expiry, blacklist)

↓ PASS

Anomaly Scorer (soft signal: "this passport's font metrics deviate from 1000 genuine samples")

↓

Risk Fusion (weighted blend)

Verdict: Add as soft signal in validation module. Rules = hard gates; ML = soft anomaly. Best of both.

2. PIPELINE ARCHITECTURE: Current vs Recommended — Decision Matrix

Current Implementation (Monolithic Sync Pipeline)

POST /scan

```
|
|
|   └─► run_ocr()           [blocking, CPU]
|
|   └─► parse_mrz()         [blocking, CPU]
|
|   └─► extract_fields()    [blocking, CPU]
|
|   └─► validate_passport/generic() [blocking, CPU]
|
|   └─► analyze_tampering()  [blocking, CPU+CV]
|
|   └─► detect_largest_face() ×2 [blocking, CPU+CV]
|
|   └─► compare_faces()      [blocking, CPU]
|
|   └─► compute_risk()       [blocking, CPU]
|
|   └─► save_scan()          [in-memory dict]
```

Problems:

- Event loop blocked for 2–10 seconds per request
- No horizontal scaling — one process = one request at a time
- No retry/dead letter — OOM or crash loses request
- No observability — no per-module latency, no distributed tracing

Architecture Doc Recommendation (Kafka + Microservices)


doc.ingested → Preprocessing → doc.preprocessed

```
|
|
|   └─► OCR Service → ocr.completed
|
|   └─► Tampering Service → tampering.completed
|
|   └─► Face Service → face.completed
|
|
|   └─► (wait for all 3) → Validation Service → validation.completed
|
|
|       └─► Risk Fusion → risk.assessed → Dashboard
```

Problems for SIH:

- Over-engineered for 48-hour hackathon
- Requires Kafka, 5+ services, service discovery, schema registry
- Debugging distributed system under time pressure = disaster
- Judges may penalize "architecture astronauting" without working demo

Recommended: Modular Monolith with Async Workers (Best of Both)



Why This Wins for SIH:

Criterion	Monolith (Current)	Kafka Micro (Doc)	Modular Monolith (Rec)
Demo Reliability	✔ Single process	✗ 5+ moving parts	✔ 2–3 processes
Scalability Demo	✗ None	✔ Real	✔ Shows worker scaling
Implementation Time Done	2–3 days	4–6 hours	
Judges Can Verify	✔ Easy	✗ Hard	✔ Easy
Production Path	Refactor	✔ Direct	✔ Incremental
Observability	Add later	Built-in	Add Redis metrics

Implementation Plan (4–6 hours):

1. Add Redis + RQ (pip install redis rq)

2. Wrap each module in a job function

```
def ocr_job(doc_bytes, doc_type): ...
```

```
def tampering_job(doc_bytes, doc_array): ...
```

```
def face_job(doc_face_bytes, live_face_bytes): ...
```

3. API route becomes:

```
@router.post("/scan")
```

```
async def scan_document(...):
```

```
    job_id = str(uuid4())
```

```
    # Enqueue parallel jobs
```

```
    q.enqueue(ocr_job, doc_bytes, doc_type, job_id=job_id+"_ocr")
```

```
    q.enqueue(tampering_job, doc_bytes, doc_array, job_id=job_id+"_tamp")
```

```
    if live_face: q.enqueue(face_job, doc_face_bytes, live_bytes, job_id=job_id+"_face")
```

```
    # Return immediately with job_id
```

```
    return {"job_id": job_id, "status": "processing", "poll":  
f"/api/documents/scan/{job_id}/status"}
```

4. Result aggregation worker (listens for job completions)

```
def aggregate_results(job_id):
```

```
    ocr_res = redis.get(f"result:{job_id}_ocr")
```

```
    tamp_res = redis.get(f"result:{job_id}_tamp")
```

```
    face_res = redis.get(f"result:{job_id}_face")
```

```
    # ... validate → risk → save
```

Frontend Change: Poll /status every 1s → show real progress (OCR done ✓, Tampering 45%...)

3. REAL-TIME FACE VERIFICATION: MediaPipe / YOLO on Camera Feed

Architecture Options

Approach	Latency	Accuracy	Complexity	Best For
MediaPipe Face Mesh	~8ms (CPU)	Good landmarks, no identity	Low	Liveness (blink, gaze), alignment
MediaPipe Face Detection (BlazeFace)	~3ms (CPU)	Fast boxes, no identity		
	Very Low	Face ROI crop for embedding		
YOLOv8-Face (Ultralytics)	~15ms (CPU) / ~3ms (GPU)	Best detection accuracy		
	Medium	Crowded scenes, multiple faces		
InsightFace (SCRFD + ArcFace)	~50ms (CPU)	Detection + Embedding unified		
	Medium	End-to-end verification		

Recommended Pipeline: MediaPipe BlazeFace (Detection) + ArcFace (Embedding) + Silent Liveness

Real-time camera feed processing (runs in browser via WebRTC or native app)

```
import mediapipe as mp
```

```
import cv2
```

```
import numpy as np
```

```
mp_face_detection = mp.solutions.face_detection.FaceDetection(  
    model_selection=1, # 0=close (2m), 1=far (5m)  
    min_detection_confidence=0.7  
)
```

```
def process_camera_frame(frame_bgr):
```

```
    """Single frame processing — call at 15-30 FPS"""
```

```
    rgb = cv2.cvtColor(frame_bgr, cv2.COLOR_BGR2RGB)
```

```
    results = mp_face_detection.process(rgb)
```

```
    faces = []
```

```
    if results.detections:
```

```

for det in results.detections:

    bbox = det.location_data.relative_bounding_box

    h, w = frame_bgr.shape[:2]

    x, y, bw, bh = int(bbox.xmin*w), int(bbox.ymin*h), int(bbox.width*w),
int(bbox.height*h)

    # Expand bbox 20% for alignment margin

    x, y, bw, bh = expand_bbox(x, y, bw, bh, 1.2, frame_bgr.shape)

    face_crop = frame_bgr[y:y+bh, x:x+bw]

    faces.append({"bbox": (x,y,bw,bh), "crop": face_crop, "score": det.score[0]})

return faces

```

Embedding extraction (run on best frame every 2-3 seconds)

```
def get_best_face_embedding(faces, arcface_model):
```

```
    if not faces: return None
```

```
    # Pick largest, highest-confidence face
```

```
    best = max(faces, key=lambda f: f["crop"].shape[0]*f["crop"].shape[1] * f["score"])
```

```
    aligned = align_face(best["crop"]) # 5-point landmarks from MediaPipe Face Mesh
```

```
    return arcface_model.get_embedding(aligned)
```

Liveness Integration (Passive — No User Action):

Silent Face Anti-Spoofing (MiniFASNet) — ONNX, runs on CPU

```
import onnxruntime as ort
```

```
liveness_session = ort.InferenceSession("minifasnet.onnx")
```

```
def check_liveness(face_crop):
```

```
    # Preprocess: resize 80x80, normalize
```

```
    inp = cv2.resize(face_crop, (80, 80)).astype(np.float32) / 255.0
```

```
inp = np.transpose(inp, (2,0,1))[None] # 1x3x80x80
score = liveness_session.run(None, {"input": inp})[0][0][1] # real class prob
return float(score)
```

Frontend Integration (WebRTC → WebSocket → Backend):

```
// frontend: capture camera, send frames via WebSocket

const stream = await navigator.mediaDevices.getUserMedia({ video: { facingMode: "user" }
});

const pc = new RTCPeerConnection();

stream.getTracks().forEach(t => pc.addTrack(t, stream));
```

```
// Send frames every 100ms (10 FPS) — backend processes, returns liveness+embedding

setInterval(() => {

  const canvas = document.createElement("canvas");

  canvas.width = video.videoWidth; canvas.height = video.videoHeight;

  canvas.getContext("2d").drawImage(video, 0, 0);

  pc.send(canvas.toDataURL("image/jpeg", 0.7));

}, 100);
```

Why Not Pure Browser Inference?

- MediaPipe JS / TensorFlow.js works but heavy (~10MB WASM), slow on mobile
- Hybrid: Browser does BlazeFace detection (light), sends crops to backend for ArcFace + Liveness
- Keeps PII (face images) on device until verification needed

4. DEEPFACE VALIDATION FOR REAL-TIME USE

What is DeepFace?

A Python wrapper library (pip install deepface) that unifies:

- Detectors: OpenCV, SSD, Dlib, MTCNN, RetinaFace, MediaPipe, YOLOv8, YuNet
- Recognizers: VGG-Face, Facenet, Facenet512, OpenFace, DeepFace, DeepID, ArcFace, Dlib, SFace
- Attributes: Age, gender, emotion, race

- Liveness: Not built-in (uses external)

Benchmark: DeepFace vs Raw InsightFace (CPU, 640×480 input)

Pipeline	Detection	Embedding	Total Latency	Memory	Accuracy (LFW)
DeepFace (RetinaFace + ArcFace)		~120ms	~80ms	~200ms	~1.2GB99.8%
DeepFace (MediaPipe + ArcFace)	~15ms	~80ms	~95ms	~800MB	99.8%
Raw InsightFace (SCRFD + ArcFace)	~25ms	~30ms	~55ms	~600MB	99.8%
MediaPipe BlazeFace + ONNX ArcFace		~8ms	~20ms	~28ms ~200MB	99.7%

DeepFace Pros/Cons for Your Use Case

Factor Assessment

Ease of Use  DeepFace.verify(img1, img2, model_name="ArcFace", detector_backend="mediapipe") — one line

Real-Time (30 FPS)  Too slow — 200ms/frame = 5 FPS max on CPU

Real-Time (5 FPS)  Possible with MediaPipe detector + GPU

Memory  Loads all models at import (TF/PyTorch) — 1GB+ RAM

Model Flexibility  Swap detector/recognizer via string args

Liveness  Not included — need separate model

Production Hardening  Not designed for servers (TF session mgmt, no batching)

Verdict: Use DeepFace for Prototyping Only

Prototype / Demo Code (acceptable for SIH demo)

```
from deepface import DeepFace
```

```
result = DeepFace.verify(  
    img1_path=document_face_path,  
    img2_path=live_face_path,  
    model_name="ArcFace",  
    detector_backend="mediapipe", # fastest CPU detector  
    enforce_detection=False,  
    distance_metric="cosine"
```

)

```
# result: {"verified": True, "distance": 0.32, "threshold": 0.68, "model": "ArcFace", ...}
```

For Production / Real-Time: Use Raw InsightFace + ONNX

```
# Production-ready: InsightFace (C++/ONNX, no TF/PyTorch runtime)
```

```
import insightface
```

```
app = FaceAnalysis(name="buffalo_l", providers=["CPUExecutionProvider"])
```

```
app.prepare(ctx_id=0, det_size=(640, 640))
```

```
# Single call: detection + alignment + embedding
```

```
faces = app.get(frame_bgr)
```

```
if faces:
```

```
    embedding = faces[0].embedding # 512-d, L2 normalized
```

```
    # Compare with stored embedding
```

```
    sim = np.dot(embedding, stored_embedding)
```

Key Insight: DeepFace is a convenience wrapper that loads heavy TF/PyTorch graphs. For real-time camera feed (15–30 FPS), you need:

1. ONNX/TensorRT models (no Python interpreter overhead in hot path)
2. Batching (process 4–8 frames together)
3. Async pipeline (capture → detect → embed → verify in separate threads)

SUMMARY: PRIORITIZED ACTION PLAN

PriorityTask	Effort	Impact
--------------	--------	--------

P0	Add SQLite persistence + basic auth token	2h	Judges see "production-ready"
----	---	----	-------------------------------

P0	Install face_recognition/dlib OR InsightFace ONNX for demo	1h	Face verification becomes credible
----	--	----	------------------------------------

P0	Remove fake progress bar → real polling	1h	UX honesty
----	---	----	------------

P1	Add PRNU module (blind, no reference needed)	4h	Detects cross-device photo swap
----	--	----	---------------------------------

P1	Add stamp template matching for visa statement	3h	Directly addresses problem
----	--	----	----------------------------

- | | | | |
|----|--|----|----------------------------------|
| P1 | Replace regex extractors with Donut (fine-tuned) | 4h | Handles all 5 doc types robustly |
| P2 | Modular monolith: Redis + RQ workers | 4h | Shows scalability architecture |
| P2 | Add Isolation Forest anomaly signal | 3h | Catches novel forgeries |
| P3 | Real-time camera feed: MediaPipe + ArcFace ONNX | 6h | Live demo "wow factor" |

Want me to elaborate on any specific innovation (PRNU implementation details, Donut fine-tuning data prep, ONNX model optimization), or proceed with a detailed implementation plan for one of the P0/P1 items?